

## 0. Project Executive Context Summary

Dimension	Description (Derived Strictly from Project Context)
Project Name	AI Co-Worker Agent Solution for Account Opening
Business Domain	Financial Services – Wealth Management Account Onboarding
Problem Being Solved	Automate and support high-volume, auditable account opening workflow for new clients, including intake, validation, custodian form assembly, review, signature, and write-back
Primary Goal(s)	Deliver AI Co-Worker agentic workflow to process ~7,000 new accounts accurately, consistently, and auditable, from intake through custodian submission and system write-back
Desired Outcome(s)	Efficient, compliant, transparent onboarding; reduced manual effort; full workflow auditability; synchronized lifecycle state across systems; robust handling of high structural variability in data
In-Scope Data Consumers	AI Co-Worker Agent Team (direct consumer), Human-in-the-loop reviewers (within workflow), Custodian platforms (data packets), Salesforce/Edge system (write-back), Subatomic Deep Lens (audit/trace)
Out-of-Scope Capabilities (Explicitly Excluded)	BI, analytics dashboards, ML training, reporting not required for agentic workflow, vector search, external data sharing, classical third-party BI/Reporting consumers
Non-Goals	General analytics/ML enablement, system modernization, business intelligence, ad hoc querying, external (out-of-ecosystem) data transmission
Key Constraints (Technical / Operational / Compliance)	All components within NorthRock IT security/infrastructure controls; auditability/transparency required at each step; must persist raw and normalized data for compliance; handle unstructured/structured data as-is; high schema and format variability; full traceability; compliance with operational SOPs; no external handling
Success Criteria	1) Full AI Co-Worker-driven orchestration of end-to-end workflow; 2) Traceable, auditable handling of all inputs/outputs; 3) Demonstrably accurate, complete custodian packet assembly and submission; 4) Write-back of lifecycle/control data to Salesforce/Edge; 5) Support for human-in-the-loop; 6) Audit validation via Subatomic Deep Lens; 7) System meets security/compliance requirements

## 1. Layer Responsibility and System Inclusion (Context-Bound)

Layer	Why This Layer Exists in This Project	Included Systems / Platforms (Data-Related; Gold may include AI Co-Workers)	Allowed Operations	Explicitly Excluded Operations	Notes
Bronze	Preserve all ingested client, custodian, and workflow (including audit) data in its received/raw form for traceability and compliance.	Salesforce/Edge, Formstack, Schwab Custodial Platform, Fidelity Custodial Platform, DocuSign, Black Diamond (optional), Client-provided Docs, Subatomic Deep Lens	Ingestion, raw object storage, schema preservation, versioning	Up-front data normalization/ transformation, field mapping, deduplication	Full format/field fidelity to source data. Handles structured, semi-, and unstructured inputs.
Silver	Normalize, classify, and contextually merge/align client intake data and supporting docs to client/household for validation and packet assembly per workflow logic.	Data Lakehouse/Warehouse platform (if required by normalization), transformation engine/process scheduler	Contextual field mapping, data normalization, validation checks, contextual linking to client/household, assembly of packet structure	Derivation not supporting agentic validation/packet build, analytics queries, BI-aggregation, inference	Directly supports Gold (Agent workflow, packet build, validation). Raw data retained for audit.
Gold	Produce shaped data packets, forms, and workflow outputs for direct consumption by AI Co-Worker Agent Team and custodian/write-back systems; enable traceable audit outputs for operational review and compliance	AI Co-Worker Agent Team (orchestrator/agents), Packet Assembly/Doc Generator, DocuSign, Salesforce/Edge (write-back interface), Subatomic Deep Lens (audit store)	Data shaping for workflow/agent events, workflow state emission, audit trace/log structuring, artifact generation, status/write-back ops	Direct reporting, BI models, anything not in workflow; external API exposure; ML model outputs	Only includes data products directly needed by agent workflow and assigned human-in-the-loop actions.

## 2. Staging Layer Design (Only If Required by Context)

Staging Area Name	Associated Layer	Included Systems / Platforms	Purpose in This Project	Data Scope	Retention Policy	Notes
Form/Doc Staging	Bronze	Temporary Object/Blob Storage	Buffer uploads and intake from Formstack, client-provided docs, DocuSign.	Unstructured PDFs, forms, envelope raw	Until load confirmation	Used to decouple high-volume/manual upload burst from persistent Bronze.
Audit Event Queue	Bronze/Silver	Event Queue/Log (as needed)	Buffer and sequence audit events and workflow traces from Subatomic Deep Lens.	Workflow trace events, audit logs	Until ingest confirmation	Optional, to ensure no loss during high-frequency workflow event ingestion.

### 3. Bronze Layer Design (Raw Ingestion Scope)

Dataset / Domain	Source(s)	Included Systems / Platforms	Storage Type	Schema Handling	Partitioning / Versioning	Lineage Requirements	Notes
Client/HH Records	Salesforce/Edge	Object Storage, RDBMS, Data Lake	Structured	As-ingested, no modification	By record ID, timestamp	Source system, extract time	System of record; supports reference and write-back.
Intake Forms	Formstack	Object Storage, Data Lake	Semi-structured	Preserve input structure	By submission/ event ID, date	Formstack, submission time	Schema varies per template; field completeness tracked.
Custodial Form Libraries	Schwab, Fidelity Custodial Platforms	Object Storage, Data Lake	Structured	As provided by custodian	By custodian, snapshot time	Custodian, form/ rule version	Differ by account type/custodian; versioning as provided.
DocuSign Envelopes	DocuSign	Object Storage, Data Lake	Semi-structured	Preserve envelope structure	By envelope/ event ID, time	DocuSign, envelope time	Includes envelope audit, routing, and artifact.
Statements (optional)	Black Diamond (optional), Client-provided Docs	Object Storage, Data Lake	Unstructured	Blob/PDF, no change	By document/ upload time	Source, upload event	Only ingested if provided/available.
Audit/Event Logs	Subatomic Deep Lens	Object Storage, Data Lake	Semi-structured	Preserve as generated	By event/trace/ timestamp	Deep Lens, originating event	Workflow/audit step traces, scoring, lineage info.

### 4. Silver Layer Design (Only If Transformation Is Required)

Dataset / Domain	Input Layer	Included Systems / Platforms (Data-Related Only)	Transformation Necessity (Required / Not Required)	Transformations Justified by Context	State Required	Notes
Intake Form Normalization	Bronze	Data Lakehouse/Warehouse, ETL/ELT engine	Required	Field mapping (schema normalization), client/household association, completeness/validation alignment per custodian rules	Yes	Enables automated validation, supports packet assembly.
Document Classification	Bronze	Data Lake, Content Classifier/Parser	Required	Document type determination, associating supporting docs to intake, basic metadata extraction	Yes	Ensures correct doc linkage for validation/packet build.
Custodian Form Alignment	Bronze	Data Lake, Rule Processing Engine	Required	Cross-reference/merge intake data with current custodian form/rule set for accurate packet assembly, validation logic checks	Yes	Required for custodian-specific packet building logic in Gold.
Audit Event Linking	Bronze	Data Lake, Audit Trace Processor	Required	Sequence/align audit events to workflow steps, enrich with metadata as needed	Yes	Ensures all workflow steps and reviews are traceable and auditable.
No other transformations	--	--	Not Required	--	--	All transformations must directly support packet build, validation, audit, or workflow requirements.

### 5. Gold Layer Design (Consumer-Driven, Context-Only)

Data Product / Output	Input Layer	Included Systems / Platforms (Data-Related; AI Co-Workers allowed if consumers)	Intended Consumer (As Stated)	Shaping Required	Refresh Strategy	Notes
Validated Account Packet	Silver	AI Co-Worker Agent Team, Packet Assembly Service, Custodial Form Generator	AI Co-Worker Agent Team, Custodian	Assemble, pre-populate per rules	Per account event	Full custodian-ready packets for submission or human-in-the-loop review; generated only on demand.
Workflow State/Exception Feed	Silver	AI Co-Worker Agent Team, Workflow State Store	AI Co-Worker Agent Team, Reviewer	Normalize exception/review data	Real-time/ event-based	Supports agent coordination, human-in-the-loop actions, and exception mgmt.
Signature Envelope Artifact	Silver	AI Co-Worker Agent Team, DocuSign	AI Co-Worker Agent Team, DocuSign	Envelope structuring	Per workflow step	Only triggered if signature required for packet; routed by AI/agent logic.

Data Product / Output	Input Layer	Included Systems / Platforms (Data-Related; AI Co-Workers allowed if consumers)	Intended Consumer (As Stated)	Shaping Required	Refresh Strategy	Notes
Audit Workflow Trace/Log	Silver	Subatomic Deep Lens Audit Store	Subatomic Deep Lens, Compliance/Review	Structured trace/event log	Continuous, step-based	For full workflow traceability, compliance, and operational review.
Lifecycle/ Status Write-Back	Silver	AI Co-Worker Agent Team, Salesforce/Edge Write-back Service	Salesforce/Edge	Structured update formatting	Per status change	Ensures lifecycle state synchronized with system of record.

## 6. Serving Stores and Specialized Layers (Only If Explicitly Needed)

Serving Layer Type	Backing Layer	Included Systems / Platforms	Why This Store Is Needed for This Project	Access Pattern	Notes
Audit/Trace Store	Gold	Subatomic Deep Lens	Ensure full traceability, audit log review, compliance monitoring	Event/Step-wise read	Only audit/trace store for operational/compliance review.
Packet/Artifact Store	Gold	Object Storage/ Content Repository	Persist generated packets/artifacts for agent/human hand-off/retrieval	Per workflow event	Stores only workflow outputs needed by agents/humans.
State/ Exception Feed	Gold	Workflow/Event State Store	Enable agent and human interaction with ongoing workflow/exceptions	Real-time/ event	Only scoped to agent/human coordination; not for BI.

## 7. End-to-End Data Flow (Context-Constrained)

Flow Step	From Layer / System	To Layer / System	Trigger / Condition	Data State	Transformations Applied	Alignment to Project Goal	Notes
1. Intake	Source Systems (Formstack, Docs, etc.)	Bronze	Submission/upload event	Raw	None (preserve as-is)	Ingest all required data, ensure original fidelity	Temp staging as needed.
2. Normalization & Classification	Bronze	Silver	On new intake event	Normalized	Field mapping, doc association, completeness checks	Prepares data for automated validation and packet assembly	Schema-flexible, traceable.
3. Validation/ Alignment	Silver	Silver	On normalization complete	Validated/ linked	Custodian/account rule validation, cross-doc consistency	Ensures packets meet submission/readiness requirements	Contextual to each account/form.
4. Packet Build (Assembly)	Silver	Gold	On validation pass	Assembled packet	Form pre-population, summary artifact creation	Enables downstream automated agent/human workflow	Generates custodian-ready packets.
5. Review/ Exception Handling	Gold	Gold	If exception/ ambiguous	Workflow state	Normalize state/exception, log action/decision	Fully supports required human-in-the-loop review/approval	Agent/human action captured.
6. Signature	Gold	DocuSign	If signature required	Envelope artifact	Envelope build/routing	Generates/routes envelopes only as needed	Only if relevant for workflow.
7. Submission/ Write-back	Gold	Salesforce/ Edge, Custodian platform	On packet completion/status	Status/ artifact	Lifecycle state update, submission artifact	Closes loop—state synchronized, data written back as per rules	Write-back via agent workflow.
8. Audit Logging	Gold	Subatomic Deep Lens	Every workflow step/action	Immutable trace	Audit, scoring, lineage event structuring	Enables traceability, operational/compliance review	Full audit at each stage.