Subatomic x Telarus - Process Discovery - October 24

VIEW RECORDING - 84 mins (No highlights)

@0:00 - **John Jungbluth (World Telecom Partners, LLC)**

We're do a quick intro. This meeting is being recorded. Walk through kind of their, what their workflow looks like today.

@0:09 - **Sam Sova**

Yeah, yeah, that was a lot of this that we wanted to get as just a deeper dive. I don't know how many of you saw it.

I just did like a quick video and sent it out of what we do. And I can certainly give maybe a high level to just level set.

And then today we really want to hear from you.

@0:28 - **Dan Foster (Telarus)**

So I think that's important, by the way, rather than us sitting there and telling you a bunch of, you know, here's our challenge.

@0:35 - **Sam Sova**

We're trying to get John faster quotes and the rest.

@0:37 - **Dan Foster (Telarus)**

We want to know what the possibilities are with you all in some of the use cases, because we take, you know, John gets an opportunity.

He thinks there's maybe a Zoom or a dial pad or a RingCentral play there. We quote it. We then help him with all of that all the way through.

He's explained that to you. But rather than us spending too much time on. And kind of here's some of the challenges.

bet you John is as adept at telling you some of that, but we really want to know what your use case and kind of how that applies to our business, if that's okay.

@1:11 - **Sam Sova**

Sure. We can start there, Dan.

@1:13 - **Dan Foster (Telarus)**

Absolutely. Does that make sense, Julie, Mari? I'm seeing that.

@1:19 - **Julie Peoples (Telarus)**

100%.

@1:19 - **Sam Sova**

Thank you. Sure. So that's great. So let's kick it off that way, and then we can go from there.

I love that. Definitely prepared to have that dialogue. So just really quickly for who's on the phone from Subatomic.

So I'm Sam Sovard, co-founder, CEO, and my background is in the enterprise. So up until three years ago, I was working in Fortune 500s just trying to solve business challenges and what the latest and greatest tech was.

I have our CTO on the phone, Carl Simon. He's in San Francisco. Carl's really built out our platform that I'll talk about.

And, and. And And had the vision for where we're at now a couple of years ago, and we started to build this.

And then we have Matt Barossi on the phone as well. Matt is our VP of sales. He's newer to the team.

I thought it would just be great for him to be on this, hear from you, and get some dialogue along the way.

So excited, you know, after walking with John through this, just really excited to talk to all of you and share what we're doing, but then also kind of hear a little bit deeper into some of the challenges you have and where AI might be able to help.

So I'm going to share a few slides here just for context, and then we can go into more of a dialogue around some of the use cases and things that we do.

SCREEN SHARING: **Sam started screen sharing - WATCH**

Can everyone see this?

@2:50 - **Mari Shankar**

Yep.

@2:50 - **Sam Sova**

Okay, perfect. So I think, I like starting here, because this is really why we built the business, right? So we are not, when we think of something,

Subatomic and I go into this, are not competing with like the Microsoft Co-Pilots or the ChatGPTs of the world necessarily.

Really the thesis we had a couple years ago looking at the market and what we thought AI would be is, you know, at some point, and we're getting to that point now, AI is really just going to be integrated into our workflows and what we do to allow us to be more efficient and drive more revenue and just have better experiences for either our partners, clients, or internally.

And these stats I show are really kind of bringing that to light. These are from the last six months from Deloitte, Gartner, MIT.

And really the core of this, you know, without just reading it directly is enterprises are really having a tough time seeing the true ROI and value of AI.

And as we dig into this, a lot of it is just based on the fact that we aren't doing what it sounds like you guys have already started to do.

You mentioned, Dan, at your offsite of just having the conversations. of like, what are some of the challenges or pain points we have and how might technology like AI really help solve those?

So what we focus on is really the complex, highly integrated, deep workflow integration of AI through what we call AI co-workers, which I'll get into now.

So if you think about, just kind of take a step back and think about all the software you've used throughout your life, really what's happened is it's a group of people sitting in a room and thinking about when we create this software for this specific instance, you know, how is it going to work?

What are the color of the buttons? Where is the, you know, where are people going to click? What's the user interface and workflow?

And it kind of fits into that box, obviously, with some customizations. And as we started to build out Subatomic, we thought, you know, AI is going to change this.

We don't... No longer do we need these silos when we think about all the tools that we have. There's an opportunity to kind of have one view of all the data that you have.

But also the way that people interact doesn't necessarily have to be in this platform that Subatomic created in this instance, right?

You can really customize it to how you want to work. So when I think about the first bullet here of integrating across systems, you'll see on the right-hand side, these are what we call our AI co-workers, AI agents is the industry term for those that are really deep into AI or following it, or maybe you just see it on your LinkedIn.

But really, this idea that we have and that we're executing on now with our platform is that we help create these AI agents, AI co-workers that can essentially first go across your systems that you want access to and kind of give that holistic view across tools or within a tool to just be more efficient.

The second thing is how you want to interact So with AI, we can create kind of any way that you want.

So what does that mean? Like a chat platform, chat GPT type platform is something pretty common that we implement, but also we're now able to interact via email, team, Slack, or even some kind of custom interface.

So what that means is like our whole premise of how we operate with our customers is AI and the way you interact with these AI co-workers we build should feel very much like how you're working with humans.

But now we're just integrating AI into it. And what that allows us to do is have this, what we call human-in-loop collaboration, right?

So we can automate a lot of things, but we can also now collaborate with AI and have contextual conversations to help do things better, faster, and more efficiently.

So our platform, very straightforward specifically, is built on this idea that we do the data piece, right? We integrate into your platforms, but then also we create what we call AI co-workers.

So this is just an example. I call it a client financial analyst. So think about this. This is the best way to think about it.

If you're going to hire somebody to do a specific task or do a role within your organization, what are you going to do?

You're going to give them, first off, when you onboard them, you're going to make sure they understand their job description.

They're going to have some sort of expertise on why you hired them. They're going to have a function within a business.

They're going to sit in an org chart somewhere. And you're going to train them. You're going to give them access to your data, your tools.

You're going to onboard them and train them for the way that you want them to work. And you're going to train them along the way.

We think of our AI co-workers as that same way. So when we start building out these specific AI co-workers for you, we look at it not as this holistic, like this is.

It's kind of the ChatGPT overarching view, right? You ask a question, you get an answer, but it's not necessarily specific to your organization or expertise within your organization.

How we think of this is we build out a team of these AI co-workers that are very specifically trained on tasks or roles that you wanted to do, and then we can pull them together into a team to get some pretty robust things done.

So a way to think about that is like this workflow, and then I'll pause and kind of gauge the room here on questions that you have, but like this is a pretty simplistic workflow on how this works, right?

We ask our question or submit kind of our task at hand that we want, or maybe it's a pre-built thing that's done.

The input is then given to the AI co-worker team. We have something called a concierge that gets every task.

So a request and then figures out the team below it to get the work done. They access your tools and systems to then do what we've trained it to do.

And they give an output. The output could be some kind of document. It could just be a next step in a workflow.

And then you get the response or kind of move on to the next phase of whatever the workflow is.

So we've created this in a way within our platform to really make this seamless for now how you work with AI.

So I'll pause there. First off, Carl, anything you want to add that I might have left out into kind of the description of this?

@9:43 - **Karl Simon (getsubatomic.ai)**

No, it was a great description. But remind one slide, I want to highlight one more thing.

@9:48 - **Sam Sova**

I'm going to annotate to do so. What do you want me to do?

@9:53 - **Karl Simon (getsubatomic.ai)**

Sorry.

@9:54 - **Sam Sova**

Oh, go back one slide, please. Oh, sure.

@10:06 - **Karl Simon (getsubatomic.ai)**

So all of this, right? Again, you onboard employees, we onboard AI co-workers, or you will, with the same information.

Of course, get a better color than yellow. How did I end up with yellow? Plus... Your tribal knowledge. Maybe you don't even have everything completely documented.

It'd be great if you did, by the way, obviously. And then it makes its way back into all this.

But along the way, you're going to continually train or bring your AI co-worker up to speed. That, you know, AI co-worker, Matt, you know, may have learned all your policies, playbooks, understands how to use your data.

And has access to all the tools it needs to do its job and has its own skill set it brings in.

But, you know, sometimes because of the way you like to do business or their best practices, they're bubbling up that haven't made its way to documentation.

You need to incorporate into that person or this AI co-worker team's understanding of how to execute for you to do it the Telarus your way.

So it makes sure it does its best practices, even if they're burgeoning or emergent. And that's through memory capture.

It learns you or we bring it up to the overall organization level where it should be a memory for everyone or reincorporate it into your cognitive way of thinking that we build behind the scenes.

Maybe we'll get into more of the tech later, depending upon who the broader audience is and the collective interest in it.

But that is a different distinction about why subatomic makes. That was it. It's a more extended statement of what was already evident in the slide, but I want to make sure you're clear.

This is what we build, and this evolves with you over time.

@12:22 - **Sam Sova**

Great. Thank you. Any questions from the group or comments that we can address?

@12:33 - **Travis Horton**

Maybe you guys will be going into this later, but regarding the tribal knowledge learning, is that happening when, you know, by observation of an employee's task completion?

Is that how that piece gets, those gaps get filled between that and the actual documentation?

@12:56 - **Karl Simon (getsubatomic.ai)**

They do. To answer that question in the best That's the way, Travis. Probably want to talk about it for at least a few minutes.

Do you want to do it now or depending on where we are?

@13:08 - **Travis Horton**

If you guys are doing that down the road, I can certainly pause. Just ask this question, open question.

@13:14 - **Sam Sova**

Yeah, Sam, what do you think in terms of? Yeah, go for it, Karl. I think it's a good thing to address.

@13:21 - **Karl Simon (getsubatomic.ai)**

There are multiple ways we do it, Travis and team. I mean, number one, it's through that natural interaction with your AI Co-Worker.

And you probably, if you've played at all in any of the model provider chat windows you've seen, at some level, if you ask for a bulleted list to be presented as a table, and you always want it for this particular item, maybe you're creating, you know, agendas for clients that you meet with, or it could be for any purpose, it will, you can actually instruct it and you can even explicitly say, hey, please.

Remember, I like to do that. And Subatomic will remember. Now, that's the easy, simple stuff. Now, we have three types of memory where we're constantly either synchronously taking your instruction saying, remember, or, and we have asynchronous where every 20 minutes, we're evaluating the nature of the discussion that has transpired, and we're picking up your semantic language, the terminology you like to use.

The episodic type of memory, where what happened in this particular interaction on chat? What did we learn about when we think about, you know, classifying the type of discussion, the types of topics, or anything that would combine in your ontology of the way you do business, in your business process, that we would learn.

And we won't want to capture it in terms of, you know, how to consider reacting in those episodes in the future.

And procedural logic. If you like things to be executed in a certain way that may be supplemental to the way you should execute a process according to your playbook, it will actually capture that as well.

And you have full administrative access to your memories also so that it's not just behind the scenes working for you.

You can actually review your memories and make edits to it directly. And then, as I alluded to earlier, promote it for overall org-level approval if you have dedicated or assigned people that can make that decision on behalf of your org, your brand, your department.

We also break it down by use case. Now, we'll talk about some of your use cases, but we have a way of segmenting the memory so that, you know, you don't have that situation where someone told you something.

About Kim Kardashian, and you're like, I didn't need to know that, now I forgot my mom's phone number. You don't need to be cluttered and have brain rot.

We have a way of segmenting the memory so that it's really tight in terms of being pulled directly for the discussion points or task execution at hand.

Now, we've talked about just memory right now, but context adaptation as an overall understanding of how things evolve with you.

To pick up on that travel knowledge. It's also in what we call the capture of your cognitive thought, your tree of thoughts as you think through things in more of its own graph-based ontology of, you know, when these topics come up, I like to think this way, you know.

And we get into really branched logic that we capture through a knowledge graph. And that is evolving and constantly updated from the...

Asynchronous processes that update memory and then move to be brought into updating, you know, the knowledge graph that has a better understanding of your business in the first place.

So much more sophisticated knowledge. And then you step back now and think about everything I've talked about. These co-workers that are implemented or defined and trained according to your needs to execute certain tasks within your workflows, then have this cognitive engine, your brain, through a knowledge graph inserted or really dynamically loaded at runtime.

Hopefully some of the language I've used wasn't necessarily too technical. tried for it not to be technical, but more understanding that there are components in our data and AI infrastructure that really allow you to make it your way, not just for today.

into the future.

@18:04 - **Travis Horton**

Thanks, Karl. Appreciate that. Absolutely. I have a general understanding, a very good understanding of what you're explaining. I'm just trying to understand we have, you know, there's a lot of complexities and, you know, we have a lot of documentation, but there's a lot of, like, little details that are probably tribal that we would need to capture.

That's the reasoning behind my question.

@18:29 - **Karl Simon (getsubatomic.ai)**

Fantastic. You know, interestingly, Travis, for one use case we did with a client in which, you know, we may get in, we may refer to it as a success story because it's in a complex space, highly regulatory, or highly regulated, and really high tribal knowledge, where we conducted this discovery, and it was deep and detailed.

And we thought, wow, okay, we captured a lot. And so it really became the basis on how we auto, you know,

Otto drove our design and engineering execution to build this thing. But there were so many little things along the way that we found out weren't originally communicated because even they forgot.

They didn't remember what they knew. And eventually, as we worked through some executions, we captured it all. And actually, let me redact or retract that last statement.

I'm sure that we'll be continuing to learn because that's the idea. This thing continuously learns, you know, but we are finally, we believe, caught all those things they couldn't remember that they do in the moment of time for executing their business process.

@19:45 - **Josh Lupresto**

Hey, would love to, as part of, if you're getting into this next, just understand those use cases a little bit.

But I would like to understand tech stack, right, where the components are, what the automation is, what you can, what can go in, what can go out, integrations, stuff like that.

Just visualize how it might fit in our world.

@20:02 - **Karl Simon (getsubatomic.ai)**

Yeah, that sounds good. We may have some sample slides that we can pull up, Sam. Yeah. it's important to your questions, Josh.

Otherwise, I can easily quickly give you a sense of, depending on what the use case was, what the text was.

@20:19 - **Sam Sova**

that's a great question. Maybe what we do, let me move to a use case, and then we can talk about, I think there's one that specifically I talk about that I can pull up a visual that might be helpful.

Awesome. So let's go on to that. Oh, no, I can't get it.

@20:35 - **Karl Simon (getsubatomic.ai)**

I got to get rid of your drawing card.

@20:38 - **Sam Sova**

let me see if I can. it. I got it, man. All right. So here's the cup. Oh, no. Here we go.

Okay. So one of the use cases, and I'm going to show a variety of industries and use cases, because I think it shows the breadth of what we can do here.

Because if you think about, you know. What we just dove into of like AI Co-Workers, similar to how you think of hiring new team members, the breadth of what we can do is pretty wide.

So this first one is just a really good use case that we hear a lot of within organizations around legal contracting.

But I think kind of the idea of what I'm going to go through of what was created can really transfer to some of the things that we already heard from John for you.

And this whole thing was, you know, Fortune 500 company, global, took three to five weeks for internal counsel to redline a contract.

As a salesperson, if there's any salespeople in the room, that's like three to five weeks is like three to five months, right?

So we helped them narrow this process down to literally three to five minutes. And how do we do that?

So similar to the process I just went through, we learned about their best practices, policies, procedures on contracts and various contracts.

contract types. And we trained a team of AI co-workers based on components of the contract. So we actually had like a warranty specialist, an IP specialist, indemnity specialist, that when these contracts came in, they knew that piece and exactly how this organization wanted to do it.

And they would redline the contract accordingly. So the process was pretty simple. The salesperson would go in, we created a custom interface for them, single sign-on.

They filled in a couple pieces of information, uploaded the contract, and literally within three to five minutes, an email would be sent to the appropriate internal counsel, internal individual.

And they would now have a redline contract that they could go through. And I mentioned human in the loop.

They can go through as a human and now accept, reject changes, and move on in the process. And to Carl's point...

And like memory and so on, like, let's say there are 10 redlines and nine were accepted and one wasn't, the system would then learn, okay, this one wasn't accepted.

Moving forward, I need to reevaluate my redlining to be more precise, right? So it's continually learning to be more and more accurate.

Another use case, the manufacturing company really focused on service, so their challenge was, you know, they would go on site to do the servicing of equipment that they've created, highly customized over the last 25 years, and they would go on site.

And there were kind of two challenges. One, the diagnosis took a lot of time because every piece of equipment was different.

So they would have to, like, go on their phone, zoom in, like, go to manuals, zoom in, and try to troubleshoot.

So we created a platform for them where... They could collaborate with their AI co-workers to diagnose an issue faster.

Then after the issue was diagnosed, they uploaded some photos, added some additional notes, and they could create an automated service report that went to the service manager for review and submittal.

Now, this is really important, not only from like an efficiency customer experience standpoint and just the ability to do more technician calls throughout the year because they're saving so much time, but also they're getting cash faster.

So, to be able to invoice the customer, they had to send these service reports, and in some cases, they were taking 60 or 90 days to send.

Now, they're literally taking a day or two to get these service reports after their internal review, which was highly automated and collaborative with their AI co-workers.

The last one I'll show you guys, and this is where I can pop up. of a visual real quick on, like, the workflow, and I think Josh just will help, like, kind of bring it to life and answer the question that you just had, is this idea of, like, wealth managers, right?

So wealth managers, they spend a lot of time not talking to clients. So about 80% of their time is spent on, like, back office type things.

And we found that this agenda creation process, which is basically preparing for client meetings and having a three-page of, like, here's your 401k, here's, you know, what we talked about last time, here's our agenda and things we want to do this time, it was taking about four hours per client a year.

They have 2,000 clients. So 8,000 hours a year it was taking just for this process. So we were able to integrate with 10 of their tools, so, like, email, custodial systems, and stuff.

And be able to completely automate this process with AI, where ahead of the meeting, the AI Co-Workers are actually going into their calendar, seeing when they have this meeting with an individual, and automatically creating this agenda.

So the advisor can now spend time looking at that, modifying it as they see fit, and spend more time focused on the in-person activities that really drive revenue for the business.

So we will save, like I said, 8,000 hours just on this process. But the big picture is they now have one view of the customer.

So what they're able to do now is essentially go into the platform, and they chose a chat platform, and ask any question about a customer at any time.

So this is super powerful, right? Like, even going into a meeting, let's say, like, oh, man, I forgot the kids' names of this client.

And they can go into the platform real quick and say, what are Sam's kids' names? I'm about to walk into a meeting, and the system will spit out, what

From all the data it has, what the kids' names are. So some really powerful stuff to even personalize it even more.

What I would like to show is just that workflow. And then I think Josh, this will answer your question.

Give me a second here. All right. So this is kind of taking everything we talked about in that use case and putting it into practice, right?

So on the left-hand side, these are all the tools that they had that we integrated to. So what does integration really mean?

What this means is we're able, through a lot of different capabilities that Karl and his team have built out, get a complete view of the customer and all the data from these platforms and put it into a data warehouse and a vector store.

And what that entails is now... All that one view of the client or the access to all client data, where on the right-hand side, we have this co-worker team that we built out that whenever a question comes in from an advisor or anyone on the team, the concierge at the top here will go in, figure out the right AI co-workers or the team of co-workers to address the question, go into the database and get an answer.

So, at a high level, that is how this works, but Josh, I'm assuming...

@28:32 - **Josh Lupresto**

I don't know if everybody does, but I'm a nerd, I do, and there's like at least one or two more nerds on here, maybe a couple in hiding.

@28:39 - **Karl Simon (getsubatomic.ai)**

Nerds Unite!

@28:41 - **Josh Lupresto**

Yeah, so, I mean, hence, I'm in Hoodie, right? So, hey, I gotta play the part. So, is that fair to say, okay, the things that they have that are data and tools have complicated or little to no APIs, maybe some screen scraping, some APIs, and then like a universal data warehouse that you can then layer on whatever front end?

And they want to integrate and chat and talk with it however they want. Is that fair?

@29:06 - **Karl Simon (getsubatomic.ai)**

you for answering your own question, Josh. That was really good. No, seriously. Yes to pretty much what you said in 30 seconds.

I mean, have, we, so depending upon the source system, you have a variety of possible ways of accessing the data.

A lot, or mostly dependent upon what they want to allow you to see, provide APIs that are partially only valid in terms of the full scope of data you want.

Or you may get it delivered through data files, through SFTP, and for pickup. But we web scrape, I mean, we use the vision, we have our own vision model operators to go in, log in, the same way you would with your license, and go ahead and pull the data.

where it can only be pulled from directly within the application. And we have those agents that are built to know how to do that scraping, to populate the database, again, based upon the knowledge of where data should land in the database, and then otherwise processes along the line.

And you also talked about, I think I heard you in that very quick statement talk about, you know, in the database landing it into maybe a more of a structured, transformed state, like it could be like a data warehouse, it could be overall, we build in a data fabric that is more federated.

But the bottom line is, we, you know, we pull from sources, quite often land it into a ready state like that, or on demand, depending upon the nature of how data can be retrieved, how timely or real time the data needs to be in support of whatever that use case is.

It's, you know, we'll go ahead and, you know, kick off, you know, where data gets pulled, and it could be multiple systems that's built into a backend workflow.

Now, the technologies, again, reading APIs, and for those companies that have built MCP servers for their systems, and more and more come online every day, we utilize those whenever possible.

As we also consider MCP conversion of human text into the actual technical pull, and conversion of your human request into a pull on the backend.

We use NADN as well for some integrations, as well as for some general dynamic workflow generation when we build out your use case.

And a lot of the agent factory that we build also, that if we don't already have an agent for it, we need to build it dynamically.

We use from invasion factory, we've, we've, we've, we've built Also, N8N and some of our own internal data infrastructure to basically say how do we want to implement or generate an agent specific for this, and we will pull across a number of different platforms that could include how we like to engineer typically, which is Lanchain, you know, but sometimes we'll Lama Index, Crew AI, or other forms, depending upon which ones are our factory.

@32:35 - **Sam Sova**

So, what I pulled up, I'm just going to scroll through this, but Josh, you alluded to, like, the platform, right?

This is, like, how they wanted it built, so we customize, kind of, the interface for the client experience that they all have access to.

@32:52 - **Dan Foster (Telarus)**

Hey, what? No, I'm just wondering, what has John shared with you on some of The opportunity and like, what's your applicability to some of that opportunity or John, have you, have you got at that level with these guys to talk about rather than interviewing us on the phone in terms of the, you know, price quoting was one of the topics that we talked about.

I mean, has, Sam, has he shared enough with you guys for you to say, oh yeah, that it's highly applicable use case that we can kind of come in and, you know, work with Mari's team, work with Julie's team to kind of deliver a solution that has ingestion, APIs, et cetera, or are you not at that level yet?

@33:36 - **Sam Sova**

So, so John, if you're okay, I'll, I'll kind of, John shared the, the high level, right? And his experience with it, what we wanted to do, especially in the second half of today was just go a little bit deeper with you guys.

So we fully understand it, but at a high level, this is, this is definitely something we can do. And hopefully, you know, with the very kind of quick overview we gave of you, you can see the different applications and, and the versatility.

But we'd love to go, I mean, we really wanted to spend time with you to go deeper into those use cases and applications.

And even if you can do like a quick show and tell, if you're okay with that, that'll give us a better understanding of it so we can ask the right questions.

@34:21 - **Ciera Broberg (Telarus)**

So I have a list of questions just after watching the video and then just watching here. And I know Travis also, we got together and we're starting to, you know, list everything that we have.

I think a lot of our questions kind of stem through a lot of different things, but the complexity with certain suppliers.

So training the AI co-worker on specific nuances like, you know, AT&T, we have to log in between six and eight systems just to check your serviceability.

So it's logging into every system. So getting, you know, pulling the serviceability, and a lot of times they're opposite.

They're not, they're never the same across the board. And then on top of that, we have the multi-factor authentication that are required for a lot of our suppliers.

So is that something that AI Co-Worker can work around as well as it's going in and out of all these different systems to come back and be like, three out of the eight systems were this serviceability, and the other ones were this serviceability?

Does it have that level of granularity from a training perspective, I guess?

@35:39 - **Karl Simon (getsubatomic.ai)**

Yeah, it does. And maybe if we answer a few questions that I might have about that, I'll give you the 100% full answer.

But first of all, yes, MFA, we built something through Twilio that we can actually go in and be able to enable MFA-based logins.

That's foundational. TableState for having AI agents be able to log in on your behalf. The second thing is for handling variability.

So it sounds like serviceability that you pull out is never the same. Do you mind giving me a little more understanding of what that means?

And then I'll tell you exactly if and how we can handle it. I know it's going to be a yes, we can.

@36:23 - **Ciera Broberg (Telarus)**

Yeah, I mean, it's a simple task. It's really just logging in, putting in the address, and it spits back availability.

This is what's available, this is what's not. We know, again, tribal knowledge kind of a thing, we know that that system isn't the end-all, be-all.

So we have other systems that we cross-check the serviceability back and basically say, oh, you know, 50% of the locations say yes, so we're just going to assume yes.

As we're, you know, going back and delivering that to the advisors, because it's just the complexity of working with that particular.

Supplier. So if it can do and make those decisions based on the outcome of the different systems.

@37:09 - **Karl Simon (getsubatomic.ai)**

Yeah. So let me, if I understood it correctly, CRM, please course correct me if I have. So first of all, yes, log in.

Whatever time restricted boundaries are required using MFA to go to specific areas, fill out fields and pull out serviceability.

You know, that may very well, or sounds like it's location based and can do it across any number of logins requires to get the full locations of, you know, where these suppliers exist.

And then it can take it and based upon your rules, your how to determine how, just the way you would today, you know, sit through the data and determine what the output or result should be.

For this given client, given their location, you know, the best. We definitely have serviceability within, maybe it's a, maybe describing a very specific, you know, certain mile radius, whatever it is, you know, you can fulfill it, and it's at the price point that you want to fulfill it.

I don't know what other factors go into it, but that is your cognitive knowledge that gets embedded into the how-to that's, you know, onboarded on day one.

And, again, if there's even travel knowledge you use in your day-to-day, it can evolve with you as you, you know, provide your feedback to your employees to help them understand a little better how to execute.

@38:37 - **Joshua Stewart (Telarus)**

Those are some really great questions, Ciera. From a use case perspective, I can kind of give you some ideas of what we're looking for in operations, right?

We're making a lot of progress in a lot of areas where we're looking for the agentic pieces, I would say, are in a few main areas.

One, of course, Ciera alluded to, which is speeding up our pricing. We work in Salesforce. We have a lot of different...

suppliers, some have APIs, some have rate cards, some provide answers through email. So we're wanting to speed up that process.

So I'll say at a high level, that would be our first use case. Second would be excessive handoffs. Even cross-functionally within our departments, just because of some of the complexity of our industry, right now we might have a few too many, but we're looking to reduce the number of handoffs between departments.

And the back and forth and identifying the ones that really aren't necessary. I would say number three, data hygiene.

And data hygiene in terms of looking for agentic for validation. For example, I'm working in Salesforce, I'm processing an order.

Perhaps I didn't enter all of the information that I know will be required to provide a quicker response from the supplier.

You know, we're gathering, hey, what's required for AT&T, for example. example, What could we use agentic for to kind of identify those and alert?

Another one would be data intelligence and analysis, right? Everywhere from performance monitoring, but mostly trend identification. We've got a lot of reports, a lot of dashboards, a lot of all that.

We're looking for agentic to see where can it be used to identify trends, right? Identify those proactive opportunities to jump in and provide better service.

Another thing is performance monitoring, right? How are we doing? Is there, and we have the dashboards and all of that, is there an agentic component that could assist there?

Another one would be escalations. We take in a lot of escalations. I would love to see if there's something that can possibly help with the root cause analysis portion of it, right?

We're collecting the information. We understand what's wrong, what happened. We want to make sure that we're really going through a proper RCA analysis.

So I think an agentic there would be great. Another one, quality assurance, right? Not only are we conducting it on time, but more so something that's really going through and scanning for customer analysis or sentiment, making sure that we understand if there's some sort of trigger or some sort of issue that maybe is just embedded in the notes and we want to surface those things.

So certainly quality assurance. Another one, of course, is communications. We know we're doing a lot with the generative stuff, but what can we do to use agentic to get to that inbox zero goal that we're all looking for?

Not only inbox zero, but are we sure that there's a direct correlation to the inbox cases and or Salesforce cases?

Some teams work through Through emails, some work through Salesforce cases. The ones that work specifically through Salesforce cases, they're still receiving emails from time to time.

I want to be sure that we have something that can really scan, go through, highlight, if possible, do a cross-reference check with Salesforce within a particular case type.

Identify if there is a case, if not, create one kind of going through that. Is there an agentic opportunity to, if we were to put a case in a pending status, for example, and kind of have it taken out of our, right now we're looking for queue, and this is, I'll give you a specific example.

So, an advisor reaches out to us, we provide an answer. Before we close that case, we want to make sure that that answer satisfies the advisor.

Sometimes it takes them a while to get back to us. So to get it out of our direct queue, right?

Imagine putting it into some pending queue. Could we have something that's just scanning that queue? And if there's anything, any response, that one, it'll surface that response.

Two, it will automatically engage possibly the advisor. It's been 30 days. You haven't responded. Are we okay to close the case?

And if they respond yes, maybe possibly automating that process. I would say FAQs, both internally and externally. And I've seen the demo, so I know that you guys can do this, but really we're wanting to build that self-help knowledge between internal and external.

So scanning our data, scanning our tribal knowledge and all of that to build a really healthy repository where our team members internally can just type the questions that we're all typing in our internal communication.

We use something kind of like Slack. Slack, it's called Glipp through, through. Central. A lot of questions and answers in there.

Hundreds, thousands. Sometimes we're asking a question that's already been asked several times. You just have to scan to the top.

Something to surface those answers I think would be great. And then finally, that I could think of, operational excellence.

Is there an agentic functionality that will scan our existing processes? I know we would have to have all those documented and outlined to automatically identify any gaps or opportunities for efficiency.

So I think that's a quick list from operations and kind of some of the things that we're looking for.

Hunter, anything to add there?

@44:44 - **Hunter Moore (Telarus)**

the only thing that I would add is when we talk about pricing, that does include the quoting process and the development and delivery of contracts.

And some of these suppliers are very particular about what we would put in a contract. We often have, you know.

Changes that may not be communicated to us from the supplier based off of what is needed in a contract.

So one of the things that I'm interested in, and I'll pass it over to my experts, but understanding what the governance of this AI looks like, especially for, you know, small changes like this very small thing now equals true.

And forever and ever, if I do this very specific type of addendum for this very specific supplier on a very specific product, now enact it.

What does that look like to train, et cetera? But that's all that I would add.

@45:42 - **Karl Simon (getsubatomic.ai)**

Great. Sounds like you only got halfway through your list because I only heard maybe double-digit list of items. So that's awesome.

I love the way you guys are thinking. I'm going to, first of all, say yes to it. Thank We, you every single one of those is in our wheelhouse.

And the second thing I want to suggest is if we wanted to go into a little more depth of why it would be, you know, in terms of the time we have available in this meeting, you know, it would be quick hitting as we walk through each one that you go by Josh or Joshua?

@46:21 - **Joshua Stewart (Telarus)**

Both. Depends on whether I'm in trouble or not. Right now, we can stick with Josh.

@46:26 - **Karl Simon (getsubatomic.ai)**

All right, Josh. Awesome. You know, for the list that you, you, you enumerated and Hunter, you added that last one for, you know, terms that should be very specific and applicable to a given supplier agreement.

While we walk through each one, and I'm going to tell you a few things. I'm going to say, number one, how to generally think about it from an overall implementation perspective.

And secondly, why subatomic in particular does this right where others may fail if it's, if it's something that is really AI.

AI-specific above and beyond typical technology automation. Sound good?

@47:07 - **Joshua Stewart (Telarus)**

Mm-hmm.

@47:09 - **Karl Simon (getsubatomic.ai)**

All right. Because I usually rely on my Fathom note-taker to take that list and tell me what we discussed, I wasn't typing while I was listening.

was focused just on listening to you intently, Josh. Do me a favor. Start on item number one. You don't have to recap what it was.

@47:27 - **Joshua Stewart (Telarus)**

It's just, I will remember what it was, and then I'll give you the response. Increasing pricing. Increasing the quoting process to provide better pricing, quicker pricing, accurate pricing through a myriad of suppliers that we work with.

Some with APIs, some with emails, some with their own portals.

@47:51 - **Karl Simon (getsubatomic.ai)**

Got it. Okay, so on the surface, this is a multi-system with complex, you know. one. go. Price to Quote, or a quote, you know, quoting system that needs to consider a series of services that may or may not be itemized or listed or maybe bundled in different ways to auto-generate the quote.

And you may have multiple suppliers that you are sourcing at the same time, right, that needs to be determined the appropriate price, even after you've applied Sierra's, you know, idea that has to be location-based within a serviceability.

Context or geography, right? So you're talking about quickly being able to integrate into systems and have an engine that is a pricing engine.

All of that can be quick. Now, I heard other things like handoffs. And by the way, the other thing I'm going to do is when I hear some of the other listed use cases, I'm going to tie it all together.

But if you identify which handoffs can be eliminated, great. If certain handoffs take time, I'd like, I would need to hear more about, you know, what.

It's causing the time. But generally speaking, you know, for the first use case, yeah, I mean, to be able to price quickly, it's always a function of being able to go into that system and gain the information.

The AI Co-Workers will be fast. Did I understand that one correctly in terms of the speed?

@49:22 - **Mari Shankar**

Okay. Sorry, Karl. With the pricing, saw Josh was, Josh is not a drummer. So, Josh, Josh was mentioning, you're absolutely right, it is multiple systems, some of them have integration, some of them have screen data, know, portals.

And as Sierra was pointing out, we go through, like, AT&T has six or eight systems where we have to log in.

It's kind of like you have to guess, right, at that point, like, once you do the serviceability check, you're making a guess, like, saying that, hey, four of these, or three of these return the same answer, so I'm going to go with that, we'll just ignore the remaining four, right?

@50:01 - **Karl Simon (getsubatomic.ai)**

Right.

@50:03 - **Mari Shankar**

So is that something that is something that we can set it up as an agent like running for AT&T and how do you manage that late age differentiation between like some sort of a logic that you will have to build it?

And the other second part of it is this can change. When we talk about these systems, maybe they'll consolidate it into four systems tomorrow or maybe they'll go to 10 systems tomorrow.

And we will come to know it only usually, and correct me if I'm wrong, Josh or Hunter, that after the fact.

Many times we will know only after that happened that they have expanded or it has been split into multiple systems that we have to log in into one more portal.

So I saw the example you were walking through Vantage, which is similar to what we are talking about from a pricing perspective, right?

It is the same methodology. We have APIs, we have portals, sometimes. Sometimes phone calls. How do you manage that?

Those kind of changes happening frequently.

@51:10 - **Karl Simon (getsubatomic.ai)**

That's a really great associated question, Mari. So let me make sure I understand when you talk about suppliers may consolidate systems or spin up new ones, and you could probably expand even that to the idea of new suppliers coming online.

How do you quickly identify what the systems are in that moment in time that are still valid, relevant for the moments of serviceability, identification, and quoting for pricing, right?

The short answer is, thanks, Dan. Really appreciate it. Obviously, this will be recorded, so you'll be able to hear everything that we cover.

So the next takeaway is, if you are at least finding out whether it's after the fact or not through email, then, yeah, we can cultivate all the sources or ways that you would be communicating today.

And just like a human being says, oh, I can't believe AT&T got rid of that system or whatever it may be, and you would react to respond.

The AI Co-Worker, have basically monitor mail and getting back to that one use case that Josh had, which is dedicated to reacting, responding to incoming communications like mail and getting to what you call the term, Josh?

@52:27 - **Joshua Stewart (Telarus)**

Inbox zero.

@52:28 - **Karl Simon (getsubatomic.ai)**

Inbox zero?

@52:30 - **Joshua Stewart (Telarus)**

stole that one from Hunter.

@52:31 - **Mari Shankar**

Zero inbox.

@52:34 - **Karl Simon (getsubatomic.ai)**

Awesome. Basically, as part of that logic, would be trying to identify supplier notifications among a number of different potential signals that you are constantly looking, you know, watching for and responding to, which would include an introduction of a new system, an elimination of another, and, you know, the, so that could be accelerated.

Now, what can or cannot necessarily be accelerated. you. you. you Is unless that supplier, when they bring up a new system, has the same, you know, accessibility instructions, like they say, go in with an API, and it's the same type of experience, if you have to log in and scrape, that would require some work to understand how to specifically connect with that new system.

That can't change unless, again, you know, until we automate the way we auto-discover differences in system integration, which is on our roadmap, but today, I'm not going to say we can spin it up today and do that.

But, you know, that's an opportunity for me to more generally state that when you subscribe with Subatomic, it is constantly adapting and delivering the best and the brightest, best in capabilities, and we are already thinking about auto-discovery on our end, on how we can.

Tackle those situations. But today, you know, the aspect of bringing that new system online would feel very similar to today, you know, pre-AI, or at least pre-subatomic for sure.

I don't think any AI can do that yet today. But at least we can shorten or compress the time period of knowing something and taking action on it.

@54:25 - **Sam Sova**

There's other thing, and I just want to double click on this. And Karl, I don't want to go too much into what I'm going to say just for the sake of time, because I want to keep hearing about these use cases.

But to Karl's point, like there's something that I didn't show in the demo just for the sake of time, but we recently spun up an audit capability where, you know, on the back end, we can always look at logs and kind of figure out how did we get to the answer that was delivered.

But now the user can do that. And one of the great things about the audit capability is you can see kind of what tools are we accessing or what data are we accessing.

It's a saying in which agentic co-workers that we've created are doing what, right? And it kind of goes through that.

And I bring this up because twofold. One, we might be able to discover some of these things in just our auditing capabilities if they change on the platforms that you use.

But second to Karl's point, like we didn't have this capability two months ago and we spun it up and now all of our clients have it as just an add-on enhanced capability.

So as the speed of AI continues to go and we spin up new things, we're just adding it to the experience of what we offer, which I think could help with a couple of these, you know, outlying things that could happen in an implementation.

ACTION ITEM: Email Sam (cc Matt) prioritized use cases + Deep Lens Qs; then schedule deep-dive/shadowing - WATCH

So I'll stop there. I want to keep hearing you talk. Karl will talk for a half an hour about auditing.

Awesome.

@55:52 - **Joshua Stewart (Telarus)**

So we can compile our list and really get it together and send it over to you so we can continue the conversation.

I know some other partners here might have some. Use Cases. Before I go, I do have a question. Sorry, Josh, I wanted to interrupt on that.

@56:08 - **Mari Shankar**

I wanted a little bit more, go on that in a deeper. Karl, Sam, this is awesome, right? What he showed on Vantage is something that we are evaluating similar stuff.

ACTION ITEM: Email Sam (cc Matt) pricing/quoting workflow; then schedule POC scoping - WATCH

Internally, we are trying to do some POCs and we have been able to make success on it. So I was hoping like next week I will have more information on the same specific use case.

We talked about the pricing. What would it take us to go into a POC mode to test it out?

What do we need from me if I want to do a POC?

@56:43 - **Sam Sova**

Yeah, we need the workflow and the details of kind of how you do it today, right? That's where we always start, the different tools that you can access and so on.

@56:54 - **Mari Shankar**

Okay. So then some of it is happening over next week for me. Sitting side by side, understanding it. So I would like to come back after that and maybe have a conversation on like, hey, how do we make a POC happen so that you can validate and test it out?

Sounds good? On the pricing piece, specifically pricing coding.

@57:20 - **Sam Sova**

Great use case.

@57:21 - **Mari Shankar**

Okay, awesome. Sorry, Josh, I wanted to get to the next steps and an action plan on that.

@57:27 - **Karl Simon (getsubatomic.ai)**

Absolutely no apology, Mari. Are we about to lose a bunch of you? Because I'm hearing, you know, people are, sounds like need to drop off.

And if so, because you have that long list, Josh, that you enumerated. Maybe others are thinking about, we'll have others I want to share.

We do need, obviously, follow-up meetings. You know, it could be of multiple nature. could be immediately dive into how we can start thinking about accelerating towards the POC.

It could be more of an extended discussion right now of what we're about to get into. Yeah. Yeah. Yeah.

But how much time do we have? I just want to ask that question first. And then if we don't have much time, we can talk about what to do next more extensively.

And I want to show you, oh, no, we lost Courtney. So I want to show you, you know, the odd, at least some screenshots of the audit.

Then it's going to answer a lot of the, how would you handle a lot of the pattern identification opportunities, the observability into what needs to be corrected, you know, how you identify things that need to be cleaned up from a hygiene perspective.

think you mentioned, Josh, identify broader patterns for insights on operational excellence. There are a number of things I can show you in four, maybe quick images to give you a sense of what our subatomic deep lens allows you for deep introspection and auditing and evaluation of what's being executed and what your systems are set up to support in this new agenda.

Authentic, you know, workflow.

@59:05 - **Hunter Moore (Telarus)**

Let's move through that demo, please. We've got about a half an hour left on the calendar.

@59:11 - **Karl Simon (getsubatomic.ai)**

And how many of us will still stay if we're a smaller group suddenly?

@59:15 - **Hunter Moore (Telarus)**

We're all here.

@59:16 - **Karl Simon (getsubatomic.ai)**

Okay.

@59:19 - **Sam Sova**

Great.

SCREEN SHARING: **Karl started screen sharing - WATCH**

@59:20 - **Karl Simon (getsubatomic.ai)**

All let's share that really fast. I'm not forgetting to wrap up any additional information we need to talk about or talk through, Josh, in terms of the first use case of the whole how to bring on new systems, pricing, and so forth.

But just so you can see a little bit more about this audit eval capability, you know, you could actually go into it.

And this is our chat interface, by the way. One of the things that we do in Subatomic overall is we also embed the capability of dynamic through dynamic react flow or react UI components, give you that, you know, rich UI experience.

Always. And check, depending upon what you're doing, right? So you can enter, you know, an evaluation. This could be for the backend for how data is processed in the first place, setting up the specific use case of the moment where you want to better understand how to perform things.

And keep in mind the whole single build out of data that allows you to do for any use case and do that pattern identification across the board.

So, sorry I don't have this an easy way to show, but you can actually go into determining even a specific system-based analysis.

It'll give you a sense of exactly how information was pulled out, but also give you indications of maybe incomplete source information that needs to be addressed.

Maybe it's information that requires hygiene or other type of cleanup. So, we actually generate a report and we generate an interactive flow.

We can move these components around. You can see different explanations and scores and retrieval metrics. Sorry, it got cut up here, but you start seeing some of the specific logic that's executed.

All of it gives you insight into how the process was executed, whether it was executed properly, and allow you feedback, a feedback mechanism.

It's, of course, correction on the fly, in this introspection of how it executed, how it performed. Here's a great example of, again, this is all one system.

We do it automatically for every single system as a part of a fast follow-up for when it actually executed a use case on your behalf.

In this case, this is agenda creation for this one client. It could be anything, but you can find out, you know, in this note for retrieving, we're evaluating.

completeness, accuracy, timeliness of the data retrieval itself, and it identifies information that's missing. So you get the full lens of really even your source systems, not just what we're executing, you know, that we built through Subatomic.

And we go through the validations that may be specific to that cognitive knowledge of, in this case, our client has clients that have age-based logic that needs, our agents need to react and respond and perform the proper level of reasoning before it knows how to actually perform a transformation.

I notice that I'm going fast, but I'm trying to be thoughtful for our 30 minutes remaining. So more on this.

Again, some, there are even specific strategies that are implemented. Your cognitive workflow, driven through not just your training, but the full knowledge graph that implements the whole intricate details of how your brains work, you know, implemented now within your ACOT.

@1:03:02 - **Joshua Stewart (Telarus)**

AI Co-Worker Team.

@1:03:03 - **Karl Simon (getsubatomic.ai)**

And you can find their findings about whether or not things are missing. this actually is downstream from the missing Medicare enrollment.

can see the cascading effect of it and how it affects Social Security claiming. And then there's a roll-up of, you know, we actually understand from the accuracy to the performance execution runtime in terms of delivering everything.

We roll everything up to have a better view of how this execution run happened. Now, keep in mind, again, this is a single execution, single system retrieval that's part of an overall more complicated workflow because you saw the list of systems on that one slide, on the left portion of the slide.

Then we roll it up to a broader level across executions. What are the patterns that you're seeing? And we identify those patterns.

And so, hopefully, as you're starting to see these examples, and this gets back, Josh, to your... In a number of these cases across pattern identification for operational improvement to source systems to identify, let alone resolve, we have everything that is fully in view and actionable.

And this becomes the input to the AI co-worker team who is going to be all about optimization of your use case processes that have been built out on how to make sure that it's improving and at minimum giving you reports on what it recommends doing if you want human and loop.

Or if you decide, I always want the agency team upon these types of findings to take action, feel the autonomy, the agency to move forward, you can give it that approval.

@1:04:52 - **Joshua Stewart (Telarus)**

Fantastic. Very impressive with the data. Quick question. I already know the answer. I just want to clarify it for the group.

One is, are these queries that you run when you're trying to do a particular task, and or is it you set it up to do this query and it just runs automatically and produces a report every time?

Is it reactive?

@1:05:20 - **Karl Simon (getsubatomic.ai)**

So you're talking about for the deep lens functionality I'm showing in particular? It's built in that it has its own, so it has its own way of, you know, building out these reports, I didn't even show an example of a report, and the flows that represent the execution path and the outcomes that were identified in the report.

So these are a part of the subatomic foundation, and it doesn't matter what your use case is, it basically auto-identifies the systems involved, the types of execution logic.

Keep in mind, what makes subatomic different is we're dynamically pulling stuff at Runtime based upon your use cases, your systems, your cognitive flow that needs to be dynamically represented in these outputs that can be visually sifted through, whether it's a flow, you know, they can interact with or report.

And the metrics and measurements based upon, again, cognitively, what you needed to do, what you believe success looked like, and generally then, you know, paired with all of our requirements of process execution, you know, whether it's performance-related latency.

If you didn't have a requirement for how quickly something comes back, we may have our own, and we may still proactively identify and recommend changes for workflow optimization that will improve the results.

@1:06:57 - **Joshua Stewart (Telarus)**

Great.

@1:06:58 - **Karl Simon (getsubatomic.ai)**

Thank you. You got it.

@1:07:03 - **Sandeep Thalluri**

That's beautiful, let me say that, Karl. That was very beautiful. Oh, yeah. I've been doing tracing, like with Trace on OpenAir platform.

Of course, I have to manually go to all of those and everything, but that was a very beautiful connected UI.

@1:07:19 - **Karl Simon (getsubatomic.ai)**

Good. Thank you, Sandeep. I appreciate that. We appreciate it. I mean, that's the important thing where it's very hard for people to do the rich, extended workflow, let alone adapt and be resilient over time to the best practices, the optimal way.

And that was important. When we built Subatomic, your way, it's not just, again, intentionally repetitive, how you want to interact with teammates, including agro workers, how you want to build in your way of thinking through things.

But it's important that you have something that helps proactively you evolve, because when you subscribe to Subatomic. Comet, you subscribe to a system that adapts at the speed of your business process or thought change.

@1:08:10 - **Sandeep Thalluri**

That's absolutely Yeah. When you say your way, that's like you guys started, like you said, two years, three years ago.

This is happening a week ago. Cloud launches its skills, cloud skills, right? Hey, it's your knowledge, your business process that happened last week.

But you guys have been doing for quite some time. You put a lot of, you know, have foresight into what's to come, what's the use cases, absolutely phenomenal.

The only question that kept through my mind was, I know how cloud skills work. I know how the chat GPT, if I want it to work, open air to work with my Salesforce, I have to dump all of my data into the vector store, all of that, and hope it makes semantic sense, right?

If I ask in our language, there is no way it would understand our vocabulary, our language, despite the fact we connect to the vector stores.

We hope it will understand. Now, is it a very similar approach that you guys have, where you would connect to our systems, which is Salesforce, it's the main system that we have, you extract all about 400 GB of regular data, plus another 20, 30 terabytes of documents we have, would you be ingesting all of that into vector store stores, and that would feed all the processes, agentic workflow processes on your system, that's how, I'm talking from a very high level, am I getting it right, or it's different completely?

If it's different, please let me know, and I'll not go deep in this call.

@1:09:41 - **Karl Simon (getsubatomic.ai)**

No, Sandeep, you asked a great question, and it makes a lot of sense, you understand it perfectly, and the answer to your question is, there is a general, you know, truth across all use cases, across all implementations, and then there's, it depends, it's a variable, you know, that, you

It's use case driven. So let me explain that a little more fully. We often think about how do we make data already ready to consume in its proper state.

And often when we extract, like you said, whether it's Salesforce as a primary source system, and they do a really good job of actually making through their APIs a lot of their data accessible, unlike a lot of others in certain domains.

We'll often both load that vector database, also a data warehouse, right? But, again, a lot of it depends on how recent does the data need to be, because we do a nightly, you know, update of both the database, of both databases, right?

Some things need to be more dynamically pulled at runtime. So often, and I used this term before, the data fabric, you know, is supportive of your whole web of intelligence, that depending upon the nature of the data that rolls into the use case, we may pull data from anything within the web of intelligence.

This includes both the downstream databases you cited or the upstream. Quite often when we think, by the way, in terms of what to store and what not to store in the vector database, it's the unstructured data or we capture a lot of our memories also that's available there if it's not already populated in cache for quick retrieval based upon the context of that workflow or that use case workflow or interaction.

You know, you time. Sorry, I could tell you were going to ask something, so I'm going to stop talking.

@1:11:39 - **Sandeep Thalluri**

No, no, no, no. I'm interested in what you have to say. The other thing that was coming to mind as an idea, probably shoot me down here, instead of we bringing data to your agents or your way of building so that we can do our way of agentic workflow, is there a possibility of bringing your agents to our data?

No. On our infrastructure, just a topic, putting it out there, that's it.

@1:12:07 - **Karl Simon (getsubatomic.ai)**

Well, I'm going to guess what you meant by it, and that means can we deploy in your cloud services?

Oh, absolutely. Okay.

@1:12:19 - **Sandeep Thalluri**

Absolutely.

@1:12:20 - **Karl Simon (getsubatomic.ai)**

It's something you should know. are, you know, when it comes to all this, we're very compliance and security minded.

And so when you, if you want us to run, you know, on our cloud services, which we support any of them, especially the top three of AWS, Azure, and Google cloud services.

But, you know, when we deploy in yours, not only are we going to be thinking about what needs to be done, but we can adopt your security requirements.

You may have an internal security team, and you have your way of doing security. So we would do it in your systems.

Again, security your way.

@1:13:01 - **Sam Sova**

Yeah. If you think about some of the examples we gave, highly regulated industries, we had to deploy within their infrastructure.

And this is where I got to give Karl a huge compliment. When we started building out the platform, this was because we didn't go into a ton of our backgrounds, but Karl and I both come from the enterprise.

Like, we're not, we haven't been startup entrepreneurs for 25 years, so we kind of knew, and Karl specifically working at very large organizations, kind of what it would take from a security standpoint to be able to do this, right?

So that was table stakes for us, to be able to do that.

@1:13:44 - **Sandeep Thalluri**

I stopped Karl, so he has a lot of Salesforce experience as a lead developer, so I stopped him.

@1:13:51 - **Karl Simon (getsubatomic.ai)**

Oh, you did?

@1:13:52 - **Sam Sova**

You looked me up on LinkedIn?

@1:13:55 - **Karl Simon (getsubatomic.ai)**

I have a long history with Salesforce.

@1:13:57 - **Mari Shankar**

We have all done that, Karl.

@1:14:01 - **Karl Simon (getsubatomic.ai)**

know. Yeah.

@1:14:02 - **Mari Shankar**

I just don't know where you live. That's the first thing we do. Awesome. No, but that's awesome that you thought about that as you started, right, deploying on our infrastructure, because that is one of the things that will be important for us, because we have complaints that we need to add a head to, force to add a head to.

So, that would require us to do some of those tweaks, depending on what data we are talking about.

@1:14:34 - **Karl Simon (getsubatomic.ai)**

Yeah.

@1:14:35 - **Sam Sova**

Yeah, and I heard one thing, too, that I want to address, Sandeep, that you hit, is I think, if I heard you right, you said, can you deploy your agents, right?

Right, so, and Carl, you can double-click on this, how you see, but there, we have a narrative, and I think, I'm going to bring this up a level, maybe for those that aren't as technical and that I sit in that, that candidly.

Um. Um. So we have a core set of agents that we've built that are pretty applicable to almost everything that we deploy.

So think of it like somebody actually said this in a meeting. were in New York this week and they said, oh, it's like the hammer and nails type people, right?

Just task oriented, just doing a bunch of different tasks and the ability, think about the ability to do screen captures and turn into data, to summarize content, to move this from here, right?

The very, the hammers and nails, right? And then we have like the subject matter expert agents that we build into that.

And that's you, right? That is your specific, like your tribal knowledge, all the things we talked about for the last hour or so.

Like that's how we think about how we build this and deploy it, right? So we have these core set of agents that are already being used by everybody, right?

That we deploy. And I just wanted to make that clear because I think, Sandeep, that's, that's something that you alluded to.

And if you didn't, hopefully that was still helpful. No. To kind of understand that.

@1:16:04 - **Sandeep Thalluri**

Definitely, our agents are our agents, because you have to be custom-specific. Now, you would have built a lot of accelerators.

An example is you take the prompt and drive the screen scraper, right, from there on, which can be probably reused across multiple use cases, because that an agent that you can bring to us, instead of we, shipping our credentials, all of that, to your infrastructure.

That was the question.

@1:16:27 - **Sam Sova**

You nailed it. Okay.

@1:16:29 - **Karl Simon (getsubatomic.ai)**

Having said that, Sandeep, you are alluding, whether intentionally or not, to the fact that it, you a lot of what we can build, these agents could be multi-use case purpose.

The idea is to have your data, the data fabric that we build for you, ready for multiple use cases.

And the agents that work within that web of intelligence, ready for multiple use cases. And that's the brilliance of being able to, when you start with one use case that might require data that...

It's shared across multiple use cases. The add-ons for tackling a lot of your efficiencies, the way we build it, allow you to move quickly and add on all these capabilities.

And you have a long list of, you know, interested or interesting use cases you want to move forward, not just one.

So between that and our interest to help further speed you along, when we talk about that adaptability and resilience, you know, we didn't talk about it, but we have built something called subatomic nucleus for our own internal capabilities, which allow us to take those hammer and nail agents, core agents, single task, single skill focused, they get rolled up into the subject matter experts, SMEs that are meta agents with a number of those eight, you know, core hammer and nail, you know, tool based or, you know, skill based agents as tools within their, you know, skill set that are then.

rolled into the workflows, your workflows, with your cognitive understanding and knowledge through the knowledge graph. But if any of those things from the bottom of that foundation are not existing, we have our agent factories, I mentioned it earlier, and through Nucleus, we can understand more about what needs to be built.

And we have our own coding agents that auto-generate the pieces that bring those agents to life fast, and so that we can actually build things out quickly, whether it's an entire new feature.

Actually, that's the one thing I didn't talk about on top of it. The workflows roll into features that are pre-built that just require your way, your cognitive approach.

But if the feature doesn't exist, we have internal development that can speed along the delivery of everything that is missing or needs to be supplemented at the hammer and nail ages all the way up to the feature generation.

@1:19:04 - **Sam Sova**

Cool.

@1:19:06 - **Karl Simon (getsubatomic.ai)**

Any other... need it to stay ahead to the Anthropics as they roll out. Things like today's skills, that's not enough, but that's...

We knew that that day would eventually come. So it's also for us as much as it is for you to be completely transparent.

@1:19:24 - **Sam Sova**

So I want to be cognizant on time here, and I thank you all for spending an hour and a half with us.

This has been fantastic. I think, you know, Josh, we have some takeaways for sure. Hunter as well, and others that shared use cases.

think kind of that next step of getting a deep dive would be good. Before we kind of wrap it up here, is there anybody else that wants to talk through any use case we might have missed or ask a question we weren't able to answer?

Are we all ready to go to whatever we're doing this weekend and talk to everybody about AI? AI and AI agents.

@1:20:06 - **Hunter Moore (Telarus)**

Tell me about your loved ones.

@1:20:08 - **Ciera Broberg (Telarus)**

I have a lot of questions, but I sent those over. think we'll just, I mean, honestly, we need our own session to go through everything.

Because I think some stuff is just really complex and unique, and we just want to make sure that's captured.

But, yeah, I think we'll coordinate that with Josh and you guys.

@1:20:25 - **Travis Horton**

Yeah, I think we'd love to get, just get some live, live shadow on, shadowing going too, just so you guys can see what we mean by complexity and what you guys, how you guys, where that translates to what you guys view as complexity or not.

@1:20:44 - **Sam Sova**

That's the best way for us to get our arms around it and give you a really good idea, too, on, like, what something like this would take from a scope perspective and so on, right?

So I love that as something you'd be willing to do. It's super helpful.

@1:20:58 - **Karl Simon (getsubatomic.ai)**

Yeah, if you, in that next step, want to think about. Ciao. Well, what would be most beneficial to you to start a deep dive?

And you're going to prioritize things in terms of your expected impact, right? What is costly today or what can actually drive revenue opportunities for you in an order of magnitude relative to today because you can move faster, more quickly priced, whatever it may be.

Most of that will be your guiding principles and coming up with a, you know, which ones should we focus on first.

keep in mind also the fact that maybe you want to group all those use cases also by the source systems that you pull from so you can leverage that one to many that you're going to have a number of those use cases.

You can look at it as almost like a bundle of use cases that could all leverage this data that would be beneficial to move on as well.

So, you know, and Josh's list had use cases. It also had really characteristics of AI that you want embedded.

like Okay. You You We'll So you can think through his long double-digit list and think, what's a feature that we just need to know subatomic acts and behaves that way and does these things for us in any use case versus true use cases?

Great point.

@1:22:22 - **Joshua Stewart (Telarus)**

Cool. Thank you. Very impressive. Anytime you were saying something I didn't understand that was too technical, I looked straight at Sandeep and Mari, and they were smiling.

So they smiled and made me smile. I knew that we were all happy on this side. So thank you.

@1:22:38 - **Sam Sova**

That's great.

@1:22:39 - **Karl Simon (getsubatomic.ai)**

Love it. Again, you know, if there's anything that we weren't clear on, definitely, I don't know if you'll run it through John or if John wants.

@1:22:51 - **John Jungbluth (World Telecom Partners, LLC)**

Yeah, I was going to say, I don't know if John wants to be the, I think John is the intro guy and now he's going to step aside.

I mean, feel free to, I guess, reply to this. in

@1:23:04 - **Sam Sova**

Yeah, why don't, anything that you guys have, just, why don't you run it through, through me and just make sure Matt is copied, if you're okay with that, and then we can, we can quarterback this.

ACTION ITEM: Confirm comms channel w/ Mari & Sandeep; update Sam (cc Matt) - WATCH

@1:23:17 - **John Jungbluth (World Telecom Partners, LLC)**

So, I actually have an idea, since, uh, Telarus does use RingCentral Clip, did you want to use that as, potentially, like, a platform to have Sam and Karl join a group there, and you can use that as the platform to share ideas?

@1:23:35 - **Hunter Moore (Telarus)**

Well, I think we can figure it out, I do need to talk with Mari and Sandeep about how they want to do it, but, uh, you can start by email, and we'll get back to y'all, certainly, sometime next week.

@1:23:45 - **Sam Sova**

Perfect. Okay. Wonderful. Well, it's so nice to meet all of you. Thank you for spending so much time with us on a Friday, especially.

Have a great weekend. Thank you, Paul.

@1:23:56 - **Matthew Buerosse**

It's very informative.

@1:23:56 - **Julie Peoples (Telarus)**

Thank you so much.

@1:23:57 - **Matthew Buerosse**

Thank you.

@1:23:59 - **Travis Horton**

Bye.